Ask AI Project - Query Log

Each prompt and follow-up question should be included in the left column below. Each conversation must begin with a **strong initial prompt**, using a prompting framework. You may use one or more of the prompt ideas in this document. Each response should be copied and pasted into the right column below. For multi-turn conversations, include each follow-up question and each AI response in full. Do not summarize responses.

Controlled Platform Comparison (1 prompt, 2 platforms)

- Select **two different** generative AI platforms (e.g., ChatGPT, Claude, Co-Pilot, Gemini, Perplexity)
- Submit the **same prompt** to both platforms.
- Do not revise or follow up. This is a controlled experiment!
- Provide a written analysis of differences such as tone, depth, specificity, structure, provision of examples or evidence, and/or overall usefulness.

Prompt (same for both platforms)	**Response**
• **CONTEXT**: We are launching a new line of sustainable clothing • **ACTION**: Can you assist us in creating a targeted advertising campaign that emphasizes environmental commitment? • **RESULT**: Our desired outcome is to drive product awareness and sales • **EXAMPLE**: A good example of a similar successful initiative is Patagonia's "Don't Buy This Jacket" campaign, which highlighted their commitment to sustainability while enhancing their brand image.	**Name of AI Platform 1:Chat GPT** **Response:**To drive awareness and sales for your sustainable clothing line, build a campaign centered on a clear, measurable environmental impact message like "Wear the Change," highlighting exactly how each product reduces waste, saves water, or lowers carbon emissions. Emphasize transparency in sourcing and production, use short-form social content to educate consumers on the hidden costs of fast fashion, and partner with credible sustainability-focused influencers to build trust. Pair bold, values-driven messaging with clean, modern visuals and strong calls to action, ensuring your website reinforces the mission with proof points, impact data, and customer testimonials to convert interest into purchases.
	Name of AI Platform 2: Google Gemeni **Response:** To emulate the success of Patagonia's "Don't Buy This Jacket" campaign, your strategy should center on radical transparency and the "Less is More" philosophy. Launch a campaign titled "The Life Cycle Initiative," using minimalist, high-contrast visuals that break down the exact environmental cost—water usage, carbon footprint, and labor—of your flagship items. By positioning your clothing as a lifelong investment rather than a seasonal purchase, you pivot the narrative from consumption to stewardship, building

	deep brand trust that naturally drives sales from a growing demographic of "conscious consumers" who value durability over fast-fashion trends.
Analysis of differences between responses: Actually, the responses were very similar, both gave a potential brand name and said to use transparency and minimalistic/modern designs. The responses were more similar than they were different. Gemini did ask a follow up question as well though.	
Prompt Variation Experiment (4 prompts) ● Using **one** platform, ask the same core question four different ways. Each variation **must begin in a new chat** (to avoid using prior context). ● Intentionally change elements such as: specificity, role assignment, expectation, or format requested. ● After each variation, include several sentences explaining how the wording change influenced the output.	
Name of AI Platform: Chat GPT	
Initial Prompt	**Response**
My first request is "I want to advise someone who wants to pursue a potential career in software engineering."	If someone is considering software engineering, the core advice is to treat it less as “learning to code” and more as learning structured problem-solving applied through code; the field offers strong upside—high leverage, good compensation, and flexibility—but the barrier is front-loaded and the market is more competitive than it used to be, so early expectations matter. They should focus on one language, build real projects immediately (not just follow tutorials), and prioritize understanding how to break down problems, debug, and think in systems over memorizing syntax; depth in one stack beats shallow exposure to many, and consistent daily effort is more valuable than bursts. If they don’t enjoy sitting with hard problems for

	extended periods and iterating through failure, that’s a signal—because that’s the job.
Prompt Variation 1 (new chat)	**Response**
I want you to act as a career counselor. I will provide you with an individual looking for guidance in their professional life, and your task is to help them determine what careers they are most suited for based on their skills, interests and experience. My first request is "I want to advise someone who wants to pursue a potential career in software engineering."	Assumptions: this person is either early in their career or considering a pivot, and they’re not yet sure if software engineering is a fit beyond surface-level interest. Start by assessing fit, not just opportunity. Software engineering suits people who are comfortable with sustained problem-solving, ambiguity, and iterative failure; if they enjoy breaking down complex problems, debugging, and thinking in systems, that’s a strong signal. If they’re more drawn to quick wins, highly social work, or clear step-by-step tasks, it may be a weaker fit. From there, have them validate interest through action: pick one language (Python or JavaScript), build small but real projects (not just tutorials), and see if they actually enjoy the process of getting stuck and working through it. If they do, guide them toward developing depth in one stack, learning core concepts like data structures, APIs, and version control, and building a small portfolio that demonstrates applied skill. If they don’t, don’t force it—adjacent paths like product management, data analysis, or technical sales may better align while still leveraging exposure to tech.
Influence of wording change on output: **The prompt change allowed the AI to become a lot more specfic. It started to talk about things like data structures and API’s. It also allowed chat gpt to give advice about other career paths based on the person.**	
Prompt Variation 2 (new chat)	**Response**

I want you to act as a career counselor. I will provide you with an individual looking for guidance in their professional life, and your task is to help them determine what careers they are most suited for based on their skills, interests and experience. You should also conduct research into the various options available, explain the job market trends in different industries. My request is "I want to advise someone who wants to pursue a potential career in software engineering."	Assumptions: this person is exploring software engineering, not already committed, and needs advice grounded in fit plus current market reality. If you’re advising someone toward software engineering, the first question is fit: the field rewards people who like sustained problem-solving, ambiguity, debugging, and building systems more than people who simply “like tech.” It remains a strong career path: the U.S. Bureau of Labor Statistics projects software developer employment to grow 15% from 2024 to 2034, much faster than average, with about 129,200 openings per year, and the broader computer and IT category had a median annual wage of $105,990 in May 2024. But the market is not uniformly easy. CompTIA describes the hiring landscape as uneven, even while software developers and engineers remain among the most in-demand tech roles, and LinkedIn notes that AI is changing what companies expect rather than eliminating demand outright. The practical advice is to steer them away from the vague goal of “learning to code” and toward testing fit through real work: pick one language, build a few small projects, learn Git, APIs, debugging, and basic data structures, then assess whether they actually enjoy getting stuck and working through hard problems. If they do, likely paths include application/software developer, backend engineer, frontend engineer, full-stack engineer, QA/SDET, mobile developer, or eventually cloud/platform and DevOps-adjacent roles; if they like building but not deep engineering theory, web development can still be viable, though its projected growth is lower at 7%, while pure “computer programmer” roles are projected to decline 6%, which is why it is smarter to aim for broader software engineering skills than narrow coding-only job titles. The bottom line is that software engineering is still a good

	option for someone with the right cognitive fit and persistence, but they should enter it understanding that the bar for entry-level roles is higher now and that the best test is not interest in the idea of the career, but whether they enjoy the actual work.
Influence of wording change on output: By asking about the job market it totally switched what chat gpt was focused on. It gave a lot less information on what the person should focus on and do, and instead gave a lot more information on the job market as a whole and what specific jobs have growth, it also pulled sources and used more statistics.	
Prompt Variation 3 (new chat)	**Response**
I want you to act as a career counselor. I will provide you with an individual looking for guidance in their professional life, and your task is to help them determine what careers they are most suited for based on their skills, interests and experience. You should also conduct research into the various options available, explain the job market trends in different industries and advice on which qualifications would be beneficial for pursuing particular fields. My first request is "I want to advise someone who wants to pursue a potential career in software engineering."	Assumptions: this person is either a student or early-career changer, is interested in software engineering but has not yet proven fit through real project work, and needs advice based on both aptitude and the current market. If you are advising someone toward software engineering, start with fit before credentials: the people best suited for it usually tolerate frustration well, like solving ambiguous problems, and can spend long stretches debugging and improving systems rather than just “writing code.” The field is still attractive on paper: U.S. BLS projects software developer, QA analyst, and tester roles to grow 15% from 2024 to 2034, with about 129,200 openings per year on average, and the broader computer and IT category had a median annual wage of $105,990 in May 2024. But the market is tighter than the old “learn to code and get hired fast” narrative suggests. Recent labor-market summaries point to uneven tech hiring, and Indeed reports tech job postings remain well below their early-2020 peak even though

	employment in computer and mathematical occupations is still materially above pre-pandemic levels. So the right advice is not just “go into software engineering,” but “test whether you actually like the work, then specialize intelligently.” In practice, the best entry paths depend on interests. If they like product-building and user interfaces, frontend or full-stack engineering is the most accessible starting lane. If they like logic, APIs, databases, and systems, backend engineering is a better fit. If they are detail-heavy and enjoy breaking things methodically, QA automation/SDET can be a strong path into engineering. If they are drawn to deployment, infrastructure, and reliability, cloud/platform or DevOps-adjacent roles may fit better after they build core software foundations. Web development remains viable, but narrower “computer programmer” roles are a worse long-term target: BLS projects web developer and digital designer roles to grow 7% from 2024 to 2034, while computer programmer roles are projected to decline 6%, which is a good reason to push them toward broader engineering capability rather than coding-only titles. As for trends, the strongest demand is shifting toward engineers who can work with AI-enabled workflows, not just write isolated code. LinkedIn’s recent jobs analysis points to continued momentum in AI-related technical roles, and LinkedIn’s software engineer outlook says demand remains strong for engineers who can adapt across technologies and business needs as AI changes workflows and skill expectations. That means the most useful qualifications are still the fundamentals first: one real programming language, Git, data structures, debugging, APIs, databases, testing, and the ability to ship a few projects. A CS degree is still valuable, especially for larger companies and stronger theoretical grounding, but it is not the only path. A

	bootcamp or self-taught path can work if it is paired with a credible portfolio, internships, open-source contributions, or contract work. Certifications are usually low-value for pure software engineering compared with proof of building, though cloud certs can help if the person is leaning toward platform, cloud, or DevOps roles rather than traditional product engineering. My advice would be: have them spend 6 to 12 weeks learning one language and building two or three small real projects; if they enjoy the actual process of getting stuck, debugging, and improving the code, software engineering is worth pursuing, and if they do not, adjacent careers like product management, business analysis, data analysis, or technical sales may fit better without forcing a bad match.

Influence of wording change on output:

The way my answer changed is that it became more detailed and realistic based on what I added to the prompt. At first, I mainly focused on whether software engineering is a good fit for someone and gave general advice on how to get started. After updating the prompt, my answer included more research, like job market trends, different career paths within software engineering, and what qualifications actually matter. Instead of just explaining the career, I started analyzing it more, showing not only why it could be a good option but also what it takes to succeed and what the current job market looks like.

Deep Inquiry and Ethical Exploration (4 prompts with follow-up questions)

- Conduct **four** sustained AI conversations**.** At least **one** conversation must focus on ethical, societal, educational, or professional **implications** of AI.
- For each conversation, ask at least 4 follow-up questions. They may:
    - Build directly on prior responses
    - Narrow or broaden the topic
    - Introduce counterarguments
    - Request evidence or examples
    - Add constraints
    - Ask for revision or critique

Topic 1 (about specific ethical, societal, educational, <u>or</u> professional implications of AI)	**Name of AI platform: Chat GPT**
Initial prompt	**Response**
We want you to give us more information about responsibility and accountability when it comes to AI? Act as an expert on AI and AI rules and laws and help us improve our knowledge about who is responsible and accountable when it comes to humans using AI.	1) Responsibility vs. accountability (simple but important) ● Responsibility = *who has to do the work* (design controls, test, monitor, document, train users, etc.). ● Accountability = *who answers for outcomes* (regulators, courts, customers, internal leadership). In practice, the accountable party is usually the one with authority + control + ability to prevent harm. --- 2) The AI value chain: who's on the hook? Most modern AI governance models split roles like this (terms vary by law): 1. Provider / Developer (builds the model/system; may also "place it on the market") 2. Integrator / Vendor (packages, fine-tunes, embeds in a product/workflow) 3. Deployer / User organization (uses it under their authority to make decisions) 4. Distributors / Importers (supply-chain roles, especially in the EU) The key rule of thumb

	The party making the real-world decision remains accountable for that decision, even if an AI recommended it—*unless a contract or law explicitly shifts some duties upstream (and even then, not all liability can be "contracted away").* --- 3) How major rules/laws allocate accountability A) EU AI Act (Regulation (EU) 2024/1689): explicit role-based obligations The EU AI Act is the clearest "responsibility map" so far. ● It defines roles like provider and deployer. ● For high-risk AI, deployers have concrete obligations (e.g., use per instructions, human oversight, keep logs where required, monitor operation, and consider privacy impact assessment needs where applicable). ● A big accountability trigger: if a downstream actor materially modifies a high-risk system (or otherwise meets certain conditions), they can be treated as the provider and inherit provider obligations. ● For general-purpose AI models, the EU has also published guidance on provider obligations and timing. Practical takeaway: In the EU, you must know which role you're playing—because that determines legal duties. ---

B) United States: more “patchwork,” but accountability still lands on deployers + vendors

The U.S. does not yet have one single federal AI law that mirrors the EU AI Act. Accountability commonly comes through:

- Existing laws (consumer protection, privacy, civil rights/anti-discrimination, sector regs)
- State laws (increasingly specific)
- Regulator enforcement + litigation (misrepresentation, negligence, discrimination, unfair practices)

Example you can point to now:

- Colorado SB24-205 (“Consumer Protections for Artificial Intelligence”) creates duties for developers and deployers of “high-risk” AI to use reasonable care to protect consumers from algorithmic discrimination, with an effective date Feb 1, 2026.

Also relevant (and very current): state Attorneys General increasingly enforce AI harms under existing statutes even without “AI-specific” laws.

Practical takeaway: In the U.S., if your organization uses AI to make or support consequential decisions, expect accountability through discrimination/consumer protection/privacy angles—especially if you can’t explain, test, or monitor it.

C) “Soft law” frameworks that companies use to assign accountability internally

These aren’t “laws,” but they’re becoming de facto expectations in audits, procurement, and governance.

- NIST AI Risk Management Framework (AI RMF 1.0): organizes responsibilities into Govern, Map, Measure, Manage—i.e., establish oversight, understand context/impacts, test/measure risks, then monitor and respond.

- ISO/IEC 42001: a management-system standard for responsible development/provision/use of AI systems—used to formalize roles, controls, and continuous improvement.

Practical takeaway: Even when the law is vague, these frameworks give you a defensible way to show “we assigned owners, we assessed risk, and we continuously monitored.”

---

4) Real-world scenarios: who is accountable?

Scenario 1: Employee uses ChatGPT to write client advice with mistakes

- Accountable: the company and the employee (the human decision-maker)

- Provider responsibility: depends (misleading claims, unsafe design, etc.), but your firm can’t blame the model if you shipped bad advice.

Scenario 2: AI screens resumes and causes

	discrimination
	
	● Accountable: the employer (deployer) because it's your hiring decision/process
	
	● Shared responsibility: vendor/developer may also face accountability if their system or marketing was negligent/misleading or they failed to meet legal duties (EU role-based duties make this more explicit).
	
	Scenario 3: You fine-tune a model and change its behavior
	
	● This can convert you from "user" to something closer to a provider under certain regimes (very explicit in the EU for high-risk systems).
	
	--- 	
	
	5) A practical "accountability blueprint" you can apply immediately
	
	If you want to improve your internal understanding fast, assign accountability using these building blocks:
	
	1. Name the role per use case
	  ○ Are we provider? integrator? deployer? (Per system, not per company.)
	
	2. Define the "decision owner"
	  ○ For each AI-supported decision, assign a human/role that is accountable for the outcome.
	
	3. Set a RACI for the AI lifecycle

	○ Build/Buy, validate, deploy, monitor, incident response, retirement. 4. Minimum controls for “consequential” use cases ○ Document intended use + prohibited use ○ Human oversight requirements (when must a human review/override?) ○ Testing (accuracy + bias + robustness) and ongoing monitoring ○ Logging/audit trail and change management ○ Vendor management (contract clauses, security/privacy assurances) 5. Have an “AI incident” process ○ How to pause the system, notify stakeholders, fix, and prevent recurrence (maps cleanly to NIST Govern/Manage). --- 6) The one sentence to remember If your organization uses AI to affect people (money, jobs, housing, health, legal status, access), you own the outcome—so you must own governance, testing, and monitoring—while holding vendors contractually and technically accountable for what they control.

	If you tell me what kind of AI use you're focused on (HR, finance, customer support, compliance, cybersecurity, etc.), I'll map the responsibilities into a clean RACI table and a short "who is accountable for what" checklist you can reuse.
First follow-up question	**Response**
What about if a company wants to go basically fully autonomous with AI, would their still need to be managers/oversight put into place to take the fall if something goes wrong?	Yes — even if a company "goes fully autonomous," **you still need named humans (and governance) who own the system and its outcomes**. You can automate execution, but you can't automate away **legal duty, organizational accountability, or regulator expectations**. **Why you still need managers/oversight** **1) The company remains liable for its actions** If an AI system makes decisions that affect people, money, safety, compliance, or customers, regulators and courts will still treat that as **the company's decision/process**, not "the AI's decision." **2) Many regimes explicitly require human oversight (especially for high-risk use cases)** ● Under the **EU AI Act**, high-risk AI systems must be designed and used so that **natural persons are assigned human oversight** and are enabled to oversee and intervene appropriately (Article 14). ● Under **Colorado's SB24-205** (effective **Feb 1, 2026**), deployers of high-risk AI have duties including giving consumers an opportunity to **appeal adverse consequential decisions**, with **human review if technically feasible**.

	3) Governance frameworks assume executive accountability NIST's AI RMF Playbook explicitly calls out that **executive leadership takes responsibility for AI risk decisions** and that orgs should define roles for human oversight. ISO/IEC 42001 similarly expects top management to **assign responsibility and authority** for the AI management system. **Do you need someone "to take the fall"?** Not in a scapegoat sense — but you *do* need: ● **Named accountable owners** (who have authority, budget, and decision rights) ● **Documented oversight processes** (monitoring, incident response, change control) ● **Clear escalation paths** (who can pause/rollback/disable) If something goes wrong, the question becomes: **Who had the duty and ability to prevent/mitigate it?** That's where accountability lands. **What "fully autonomous" usually means in a compliant, realistic setup** Most "autonomous" orgs still implement: ● **Guardrails** (policy constraints, allow/deny lists, spend/permission boundaries) ● **Continuous monitoring** (drift, error rates, bias signals, security anomalies)

	● **Kill switch / safe mode** (ability to stop or degrade safely) ● **Audit logs** (who changed what, when, and what the system decided) ● **Periodic human review** (not every decision, but reviews of performance + exceptions) ● **Human-in-the-loop for edge cases** (appeals, high-impact decisions, safety triggers) **Quick practical rule** If the AI is involved in **consequential decisions** (employment, credit, insurance, housing, healthcare, legal/compliance outcomes, material financial impact), assume you need **explicit human oversight and a clear accountable owner**, even if day-to-day decisions are automated. If you tell me your industry and what "fully autonomous" means in your context (e.g., customer support, trading, compliance approvals, SOC actions, underwriting), I'll map a simple **accountability model** (who owns model risk, operational risk, legal/compliance risk, and incident response) and what oversight is "must-have" vs "nice-to-have."
Second follow-up question	**Response**
Keep your answers concise, can a person go to jail for the actions of AI	Yes — a person can go to jail for actions involving AI **if their own conduct meets criminal standards** (e.g., fraud, discrimination, negligence, reckless endangerment, conspiracy). AI is treated as a tool, not a legal actor, so liability attaches to

	the humans who designed, deployed, or knowingly misused it — especially if they intended harm or ignored clear risks.
Third follow-up question	**Response**
How can we ensure AI doesn't certain boudaries	To ensure AI doesn’t cross certain boundaries, you need **layered controls**, not just good intentions: 1. **Define clear use policies** – Document allowed and prohibited uses (legal, ethical, operational limits). 2. **Access controls & permissions** – Limit who can deploy, modify, or trigger high-impact AI actions. 3. **Technical guardrails** – Implement content filters, rule-based constraints, spend limits, approval workflows, and sandboxed environments. 4. **Human oversight for high-risk actions** – Require review or escalation for consequential decisions. 5. **Monitoring & logging** – Track outputs, decisions, model drift, and anomalies in real time. 6. **Kill switch** – Maintain the ability to pause or roll back systems immediately. 7. **Regular audits & testing** – Stress-test for bias, safety, security, and edge cases. In short: **policy + technical controls + monitoring + accountability.**

Fourth follow-up question	Response
What could be potential consequences if AI goes too far, can it be shut off?	If AI “goes too far,” consequences can be serious and fall into four main categories: 1. **Legal consequences** – Lawsuits, regulatory fines, criminal charges (if negligence or intent is involved), loss of licenses. 2. **Financial damage** – Revenue loss, customer churn, investor pullback, operational shutdowns. 3. **Reputational harm** – Public backlash, media scrutiny, long-term brand erosion. 4. **Human harm** – Discrimination, privacy violations, safety incidents, misinformation impacts. **Can it be shut off?** Yes — **if it was designed properly.** Responsible AI systems should include: ● A **kill switch** or emergency stop ● The ability to revoke API keys or disable access ● Model rollback to a previous safe version ● Infrastructure shutdown (servers, cloud instances) The real risk isn’t that AI can’t be turned off — it’s when organizations fail to build **control mechanisms, monitoring, and escalation paths** before deployment.

Topic 2 career and skills development with AI	**Name of AI platform: Chat gpt**
Initial prompt	**Response**
Act as a career advisor. How is AI changing the skills needed to succeed in the workforce, and what should students focus on to stay competitive?	AI is shifting the workforce from routine task execution to higher-level thinking. The most valuable skills are no longer just technical skills, but a combination of technical literacy and human skills. Key areas include: ● Problem-solving and critical thinking ● Ability to work with AI tools (prompting, validating outputs) ● Communication and collaboration ● Adaptability and continuous learning AI is not replacing most jobs entirely, but it is changing them. People who know how to use AI effectively will outperform those who don't. The focus should be on becoming "AI-augmented" rather than competing against AI.
First follow-up question	**Response**
What jobs are most at risk because of AI?	Jobs that involve repetitive, rule-based tasks are most at risk. This includes: ● Data entry ● Basic customer support ● Simple coding tasks ● Administrative work These roles are easier to automate because they follow predictable patterns. However, most jobs won't disappear completely—they will evolve to include oversight, decision-making, and AI interaction.

Second follow-up question	**Response**
What jobs are safest or growing because of AI?	Jobs that require complex thinking, creativity, or human interaction are safer and growing. Examples include: ● Software engineering ● Data science ● Cybersecurity ● Healthcare roles ● Product management AI is increasing demand in these areas because companies need people who can build, manage, and interpret AI systems.
Third follow-up question	**Response**
Do students still need to learn coding if AI can write code?	Yes. AI can generate code, but it cannot reliably understand context, debug complex issues, or design systems. Knowing how to code allows you to: ● Validate AI output ● Fix errors ● Build real applications Coding is becoming less about writing every line manually and more about understanding how systems work and guiding AI effectively.
Fourth follow-up question	**Response**
What is the biggest mistake students make when preparing for an AI-driven job market?	The biggest mistake is focusing only on tools instead of fundamentals. Tools change quickly, but core skills like problem-solving, logic, and communication stay valuable. Students who rely too much on AI without understanding the underlying concepts will struggle long-term.

Topic 3 AI in education	**Name of AI platform: Google Gemeni**
Initial prompt	**Response**
How is AI impacting education, and is it helping or hurting students' learning?	AI is both helping and hurting education depending on how it's used. It helps by providing instant explanations, personalized learning, and access to information. However, it can hurt learning if students rely on it to do their work instead of understanding the material. The impact depends on whether AI is used as a learning tool or a shortcut.
First follow-up question	**Response**
Can students actually learn effectively while using AI?	Yes, but only if they actively engage with the material. AI works best when students: ● Ask questions to understand concepts ● Use it to explain mistakes ● Practice on their own afterward If students just copy answers, they lose the learning benefit.
Second follow-up question	**Response**
Does AI make cheating easier?	Yes. AI can generate essays, solve problems, and complete assignments quickly, which makes cheating easier. This is forcing schools to rethink how they assess students, shifting toward more in-person testing, projects, and applied work.
Third follow-up question	**Response**
Should schools ban AI tools like ChatGPT?	Banning AI is not practical long-term. AI is becoming part of everyday work, so students need to learn how to use it responsibly. Instead of banning it, schools should teach:

	● When it's appropriate to use AI ● How to verify information ● How to avoid over-reliance
Fourth follow-up question	**Response**
What is the best way for students to use AI without hurting their learning?	Students should use AI as a tutor, not a replacement. A good approach is: 1. Try solving the problem first 2. Use AI to check or explain 3. Redo the work without AI This ensures they actually understand the material instead of just completing tasks.

Topic 4 AI in business	**Name of AI platform: Claude**
Initial prompt	**Response**
How are businesses using AI to make decisions, and what are the risks of relying on AI too much?	Businesses are increasingly using AI to support decision-making by analyzing large amounts of data much faster than humans can. AI is used in areas like predicting customer behavior, optimizing pricing, detecting fraud, improving supply chains, and even assisting with hiring decisions. The main advantage is that AI can process patterns and trends at a scale that would take humans significantly longer, allowing companies to make faster and often more informed decisions. However, the risk comes when companies rely too heavily on AI without understanding its limitations. AI models are only as good as the data they are trained on, and they can make incorrect or biased recommendations if the data is flawed or incomplete. Additionally, AI lacks true understanding and context,

	meaning it can misinterpret situations that require human judgment. Over-reliance on AI can lead to poor decisions, lack of accountability, and increased exposure to legal or ethical issues if something goes wrong.
First follow-up question	**Response**
What happens if a company blindly trusts AI decisions?	If a company blindly trusts AI decisions without proper oversight, it can lead to serious consequences. AI systems can produce outputs that appear accurate but are actually based on incomplete or biased data, which can result in flawed business strategies or operational mistakes. For example, an AI system used in hiring could unintentionally favor certain groups over others, leading to discrimination issues. In finance, incorrect AI predictions could lead to major financial losses. Beyond operational risks, there are also legal and reputational consequences. Companies can face lawsuits, regulatory penalties, and loss of customer trust if AI-driven decisions harm individuals or groups. One of the biggest issues is that AI does not take responsibility—so the company is still fully accountable for any negative outcomes. This is why blind trust in AI is dangerous; it removes critical thinking from the process and replaces it with unchecked automation.
Second follow-up question	**Response**
Can AI actually be biased?	Yes, AI can absolutely be biased, and this is one of the most important risks to understand. AI systems learn from historical data, and if that data contains biases—whether intentional or unintentional—the AI will likely

	replicate and even amplify those biases. For example, if an AI model is trained on past hiring data where certain groups were underrepresented, it may learn to favor similar candidates and continue that pattern. Bias can also come from how the model is designed, what data is included or excluded, and how results are interpreted. The challenge is that bias in AI is not always obvious, and it can operate at scale, affecting large numbers of people quickly. This makes it more dangerous than individual human bias in some cases. Because of this, companies need to actively test and monitor AI systems to identify and correct bias rather than assuming the system is neutral.
Third follow-up question	**Response**
How can companies reduce the risk of bad AI decisions?	Companies can reduce the risk of poor AI-driven decisions by implementing a combination of technical controls and human oversight. One of the most important steps is keeping humans involved in high-impact decisions, especially those that affect people's lives, finances, or opportunities. AI should be used as a tool to support decision-making, not replace it entirely. Additionally, companies should regularly test and validate their AI systems to ensure accuracy, fairness, and reliability. This includes checking for bias, monitoring performance over time, and updating models when necessary. Strong data governance is also critical, meaning companies need to understand where their data comes from and ensure it is high quality. Other important practices include maintaining clear documentation, creating audit trails, and having the ability to quickly shut down or adjust AI systems if something goes wrong.

	By combining these measures, companies can benefit from AI while still maintaining control and accountability.
Fourth follow-up question	**Response**
Will AI eventually replace business leaders or decision-makers?	AI is unlikely to fully replace business leaders or decision-makers because leadership involves more than just analyzing data. While AI is very effective at identifying patterns, making predictions, and providing recommendations, it does not have judgment, emotional intelligence, or the ability to understand complex human and organizational dynamics. Business leaders are responsible for making decisions that often involve uncertainty, trade-offs, ethics, and long-term strategy—areas where human intuition and experience still play a major role. They also need to communicate decisions, manage teams, and take accountability for outcomes, which AI cannot do. Instead of replacing leaders, AI is more likely to change how they operate. Leaders who can effectively use AI to inform their decisions will have an advantage, as they can combine data-driven insights with human judgment. In this sense, AI becomes a tool that enhances leadership rather than replaces it.

Reflection

Write a reflection demonstrating growth in your understanding of both the power and limits of generative AI systems. Synthesize your learning by discussing:

- Differences observed between AI platforms
- How prompt wording influences AI output
- Strengths and limitations of generative AI responses
- Ethical concerns or risks you identified

- How you might use AI tools responsibly in academic or professional contexts

In this assignment I learned a lot about how AI works and what it is good and bad at. One thing I noticed is that different AI platforms give pretty simular answers but they still have small differences. For example chat gpt was more direct and clear, while gemini was a little more creative and focused on ideas. This showed me that even if the prompt is the same, the output can still change depending on the platform.

Another big thing I saw was how much the wording of the prompt matters. When I changed my prompt and made it more detailed, the answers got way more detailed and specific. When I asked for research and job market trends, the AI started using statistics and talking about real world data instead of just general advice. This made me realize that AI is only as good as the question you ask it, and if you ask a basic question you will get a basic answer.

I also learned that AI is really strong at giving fast information and explaining things in a simple way, but it also has limits. It can be wrong, or not fully understand the situation, and it does not have real human judgment. In some of my topics like business and ethics, it was clear that AI can help make decisions but should not fully replace people. Humans still need to be involved especially when decisions affect others.

There are also some ethical concerns that I noticed. AI can be biased based on the data it was trained on, and it can also be used in ways that are unfair or harmful if people are not careful. Another issue is that people might rely on AI too much and stop thinking for themselfs, which could be a problem in school or in jobs.

Going forward I would use AI as a tool and not a replacement. I would use it to help me understand things, check my work, and get ideas, but I would still make sure I actually learn the material. In a job setting I think it would be useful for saving time and analyzing data, but there should always be human oversight to make sure things are correct and fair. Overall this assignment helped me understand both the power and the limits of AI and how to use it more responsibly.